



**TO
THE
NEW™**



Intro To MongoDB


By :- Mayank



IntelliGrape

TO
THE
NEW™





mongoDB

{ name: mongo, type: DB }

- MongoDB is a document database that provides high performance, high availability, and easy scalability.
- Document Database
- Documents (objects) map nicely to programming language data types.
- Embedded documents and arrays reduce need for joins.



High Performance

- Embedding makes reads and writes fast.
- Indexes can include keys from embedded documents and arrays.

High Availability

- Replicated servers with automatic master failover.

Easy Scalability

- Automatic sharding distributes collection data across machines.
- Eventually-consistent reads can be distributed over replicated servers.



Latest Stable version of mongodb

3.4

To get started, there are six simple concepts we need to understand

- MongoDB has the same concept of a database with which you are likely already familiar (or a schema for you Oracle folks). Within a MongoDB instance you can have zero or more databases, each acting as high-level containers for everything else.
- A database can have zero or more collections . A collection shares enough in common with a traditional table that you can safely think of the two as the same thing.
- Collections are made up of zero or more documents .Again, a document can safely be thought of as a row .




- A document is made up of one or more fields , which you can probably guess are a lot like columns .
- Indexes in MongoDB function much like their RDBMS counterparts.
- When we ask MongoDB for data, it returns a cursor, which we can do things to,
 - such as counting or skipping ahead, without actually pulling down data.



Counterparts from SQL world

- Table become Collection
- Row become Document
- Column become Field

The core difference comes from the fact that relational databases define columns at the table level whereas a document-oriented database defines its fields at the document level.



The Windows Club

Collection

- A collection is the equivalent of an RDBMS table. Collections do not enforce a schema
- and is a grouping of MongoDB documents.
- For example “student” collection containing several documents for each student.



Documents

MongoDB documents are composed of field-and-value pairs and have the following structure:

```
{  
  field1: value1,  
  field2: value2,  
  field3: value3,  
  ...  
  fieldN: valueN  
}
```

NOTE: Field names are strings. Field names cannot contain null characters, dots (.) or dollar signs (\$). Each document within a collection can have its own unique set of fields .



CRUD

C

CREATE

R

READ

U

UPDATE

D

DELETE

Create Collection

```
db.createCollection(name, {<options>})
```

```
For example:
```

```
db.createCollection("demo").
```



Insert

```
db.collection.insert(document)  
For example:  
db.unicorns.insert({name: 'Aurora', gender: 'f', weight: 450})
```

NOTE: The `insert()` method has the following behaviors:

- If the collection does not exist, then the `insert()` method will create the collection.
- If the document does not specify an `_id` field, then MongoDB will add the `_id` field and assign a unique `ObjectId` for the document before inserting.
- Most drivers create an `ObjectId` and insert the `_id` field, but the mongod will create and populate the `_id` if the driver or application does not.
- If the document specifies a new field, then the `insert()` method inserts the document with the new field. This requires no changes to the data model for the collection or the existing documents.

Read

- | Find() command is a very usefull command for retreiving data on the basis
- | of some condition or selection criteria.

- | **Syntax:**

- | `db.collection.find(<criteria>, <projection>)`
- | where criteria and projection both are optional.
- | It returns a cursor to the documents that match the query
- | criteria.



The projection parameter takes a document of the following form:

```
{ field1: <boolean>, field2: <boolean> ... }
```

The <boolean> value can be any of the following:-

- 1 or true to include the field. The find() method always includes the _id field even if the field is not explicitly stated to return in the projection parameter.
- 0 or false to exclude the field.

NOTE: A projection cannot contain both include and exclude specifications, except for the exclusion of the _id field. In projections that explicitly include fields, the _id field is the only field that you can explicitly exclude.



Query for Equality

```
db.unicorns.find({gender:  
  'm'})
```



Query for Ranges

- \$gt (greater than)
- \$gte (greater than equal to)
- \$lt (less than)
- \$lte (less than equal to)
- \$ne (not equal to)

· For example:

```
db.unicorns.find({gender: {$ne: 'f'}, weight: {$gte: 701}})
```



Query Using Operators

- \$in
- \$or
- \$and
- \$exists

Update

Modifies an existing document or documents in a collection.

Syntax: `db.collection.update(query, update, options)`

Options:

1. **upsert**: If set to true, creates a new document when no document matches the query criteria.

The default value is false, which does not insert a new document when no match is found.

2. **multi**: If set to true updates multiple documents otherwise only first one is updated by default.



\$set operator

- \$set modifier is used to replace the value of a field to the specified value.
- If the field does not exist,
- the \$set operator will add the field with the specified value.

· For example:

```
· db.unicorns.update({name: 'Roooooodles'}, {$set:  
  {weight: 590}})
```



Some other useful operators

- \$inc
- \$push
- \$addToSet

Ordering

```
db.unicorns.find().sort({weight: -1})
```

NOTE: We specify the fields we want to sort on,
using 1 for ascending and -1 for descending.



Pagination

Paging results can be accomplished via the limit and skip cursor methods.

For example:

```
db.unicorns.find().sort({weight:  
-1}).limit(2).skip(1)
```



Count

```
db.unicorns.count({vampires: {$gt: 50}})
Or
db.unicorns.find({vampires: {$gt:
50}}).count()
```



Remove

- | `db.collection.remove(query, justOne)`
- | Removes documents from a collection.

- | `db.collection.drop()`
- | Called on a collection to drop it from the database.



Arrays and Embedded Documents

MongoDb turns out to be incredibly handy when dealing with many-to-one or many-to-many relationships.

Example:

```
db.employees.insert({_id: ObjectId("4d85c7039ab0fd70a117d733"),  
  name: 'Siona',  
  manager: [ObjectId ("4d85c7039ab0fd70a117d730"),  
            ObjectId("4d85c7039ab0fd70a117d732")]  
  })
```

```
db.employees.insert({_id: ObjectId("4d85c7039ab0fd70a117d734"),  
  name: 'Ghanima',  
  family: {mother: 'Chani',  
  father: 'Paul', brother: ObjectId("4d85c7039ab0fd70a117d730")}  
  })
```



Thank You!!

